

# Introduction to SQL

Select-From-Where Statements

Multirelation Queries

Subqueries

# Why SQL?

- SQL is a very-high-level language.
  - Say “what to do” rather than “how to do it.”
  - Avoid a lot of data-manipulation details needed in procedural languages like C++ or Java.
- Database management system figures out “best” way to execute query.
  - Called “query optimization.”

# Select-From-Where Statements

**SELECT** desired attributes

**FROM** one or more tables

**WHERE** condition about tuples of  
the tables

# Our Running Example

- All our SQL queries will be based on the following database schema.

- Underline indicates key attributes.

Beers(name, manf)

Bars(name, addr, license)

Drinkers(name, addr, phone)

Likes2(drinker, beer) → Likes2 !!!

Sells(bar, beer, price)

Frequents(drinker, bar)

# Example

- Using **Beers(name, manf)**, what beers are made by Anheuser-Busch?

```
SELECT name
```

```
FROM Beers
```

```
WHERE manf='Anheuser-Busch' ;
```

# Result of Query

| name     |
|----------|
| Bud      |
| Bud Lite |
| Michelob |
| . . .    |

The answer is a relation with a single attribute, name, and tuples with the name of each beer by Anheuser-Busch, such as Bud.

# Meaning of Single-Relation Query

- Begin with the relation in the FROM clause.
- Apply the selection indicated by the WHERE clause.
- Apply the extended projection indicated by the SELECT clause.

# Operational Semantics

| name | manf           |
|------|----------------|
|      |                |
| Bud  | Anheuser-Busch |
|      |                |

Tuple-variable  $t$   
loops over all  
tuples

Include  $t.name$   
in the result, if so

Check if  
Anheuser-Busch



# Operational Semantics --- General

- Think of a *tuple variable* visiting each tuple of the relation mentioned in **FROM**.
- Check if the "current" tuple satisfies the **WHERE** clause.
- If so, compute the attributes or expressions of the **SELECT** clause using the components of this tuple.

## \* In SELECT clauses

- When there is one relation in the FROM clause, \* in the SELECT clause stands for “all attributes of this relation.”
- Example: Using Beers(name, manf):

```
SELECT  *  
FROM    Beers  
WHERE   manf='Anheuser-Busch' ;
```

# Result of Query:

| name     | manf           |
|----------|----------------|
| Bud      | Anheuser-Busch |
| Bud Lite | Anheuser-Busch |
| Michelob | Anheuser-Busch |
| . . .    | . . .          |

Now, the result has each of the attributes of Beers.

# Renaming Attributes

- If you want the result to have different attribute names, use "AS <new name>" to rename an attribute.

- Example: Using Beers(name, manf):

```
SELECT name AS beer, manf
FROM Beers
WHERE manf='Anheuser-Busch';
```

# Result of Query:

| beer     | manf           |
|----------|----------------|
| Bud      | Anheuser-Busch |
| Bud Lite | Anheuser-Busch |
| Michelob | Anheuser-Busch |
| . . .    | . . .          |

# Expressions in SELECT Clauses

- Any expression that makes sense can appear as an element of a SELECT clause.
- **Example:** Using **Sells(bar, beer, price):**

```
SELECT bar, beer,  
        price*114 AS priceInYen  
FROM Sells;
```

# Result of Query

| bar   | beer   | priceInYen |
|-------|--------|------------|
| Joe's | Bud    | 285        |
| Sue's | Miller | 342        |
| ...   | ...    | ...        |

# Example: Constants as Expressions

□ Using Likes2(drinker, beer):

```
SELECT drinker,  
       'likes Bud' AS whoLikesBud  
FROM Likes2  
WHERE beer = 'Bud';
```



# Result of Query

| drinker | whoLikesBud |
|---------|-------------|
| Sally   | likes Bud   |
| Fred    | likes Bud   |
| ...     | ...         |

# Example: Information Integration

- We often build “data warehouses” from the data at many “sources.”
- Suppose each bar has its own relation `Menu(beer, price)` .
- To contribute to `Sells(bar, beer, price)` we need to query each bar and insert the name of the bar.

# Information Integration --- (2)

- For instance, at Joe's Bar we can issue the query:

```
SELECT 'Joes Bar', beer, price  
FROM Menu;
```

# Complex Conditions in WHERE Clause

- Boolean operators AND, OR, NOT.
- Comparisons =, <>, <, >, <=, >=.
- And many other operators that produce boolean-valued results.

# Example: Complex Condition

- Using `Sells(bar, beer, price)`, find the price Joe's Bar charges for Bud:

```
SELECT price
FROM Sells
WHERE bar = 'Joes Bar' AND
       beer = 'Bud';
```

# Patterns

- A condition can compare a string to a pattern by:
  - <Attribute> LIKE <pattern> or  
    <Attribute> NOT LIKE <pattern>
- *Pattern* is a quoted string with
  - % = "any string";
  - \_ = "any character."

## Example: LIKE

- Using Drinkers(name, addr, phone) find the drinkers with exchange 555:

```
SELECT name  
FROM Drinkers  
WHERE phone LIKE '%555-____';
```

# NULL Values

- Tuples in SQL relations can have NULL as a value for one or more components.
- Meaning depends on context. Two common cases:
  - *Missing value* : e.g., we know Joe's Bar has some address, but we don't know what it is.
  - *Inapplicable* : e.g., the value of attribute *spouse* for an unmarried person.



# Comparing NULL's to Values

- The logic of conditions in SQL is really 3-valued logic: TRUE, FALSE, UNKNOWN.
- Comparing any value (including NULL itself) with NULL yields UNKNOWN.
- A tuple is in a query answer iff the WHERE clause is TRUE (not FALSE or UNKNOWN).

# Three-Valued Logic

- To understand how AND, OR, and NOT work in 3-valued logic, think of **TRUE** = 1, **FALSE** = 0, and **UNKNOWN** =  $\frac{1}{2}$ .
- AND = MIN; OR = MAX, NOT( $x$ ) =  $1-x$ .
- Example:  
TRUE AND (FALSE OR NOT(UNKNOWN)) =  
MIN(1, MAX(0, (1 -  $\frac{1}{2}$  ))) =  
MIN(1, MAX(0,  $\frac{1}{2}$  )) = MIN(1,  $\frac{1}{2}$  ) =  $\frac{1}{2}$ .

# Surprising Example

□ From the following Sells relation:

| bar       | beer | price |
|-----------|------|-------|
| Joe's Bar | Bud  | NULL  |

SELECT bar

FROM Sells

WHERE price < 2.00 OR price >= 2.00;

← UNKNOWN →

← UNKNOWN →

← UNKNOWN →

# Reason: 2-Valued Laws $\neq$ 3-Valued Laws

- Some common laws, like commutativity of AND, hold in 3-valued logic.
- But not others, e.g., the *law of the excluded middle* :  $p \text{ OR } \text{NOT } p = \text{TRUE}$ .
  - When  $p = \text{UNKNOWN}$ , the left side is  $\text{MAX}( \frac{1}{2}, (1 - \frac{1}{2}) ) = \frac{1}{2} \neq 1$ .

# Multirelation Queries

- Interesting queries often combine data from more than one relation.
- We can address several relations in one query by listing them all in the FROM clause.
- Distinguish attributes of the same name by "<relation>.<attribute>".

# Example: Joining Two Relations

- Using relations `Likes2(drinker, beer)` and `Frequents(drinker, bar)`, find the beers liked by at least one person who frequents Joe's Bar.

```
SELECT beer
FROM Likes2, Frequents
WHERE bar = 'Joes Bar' AND
      Frequents.drinker =
        Likes2.drinker;
```

# Formal Semantics

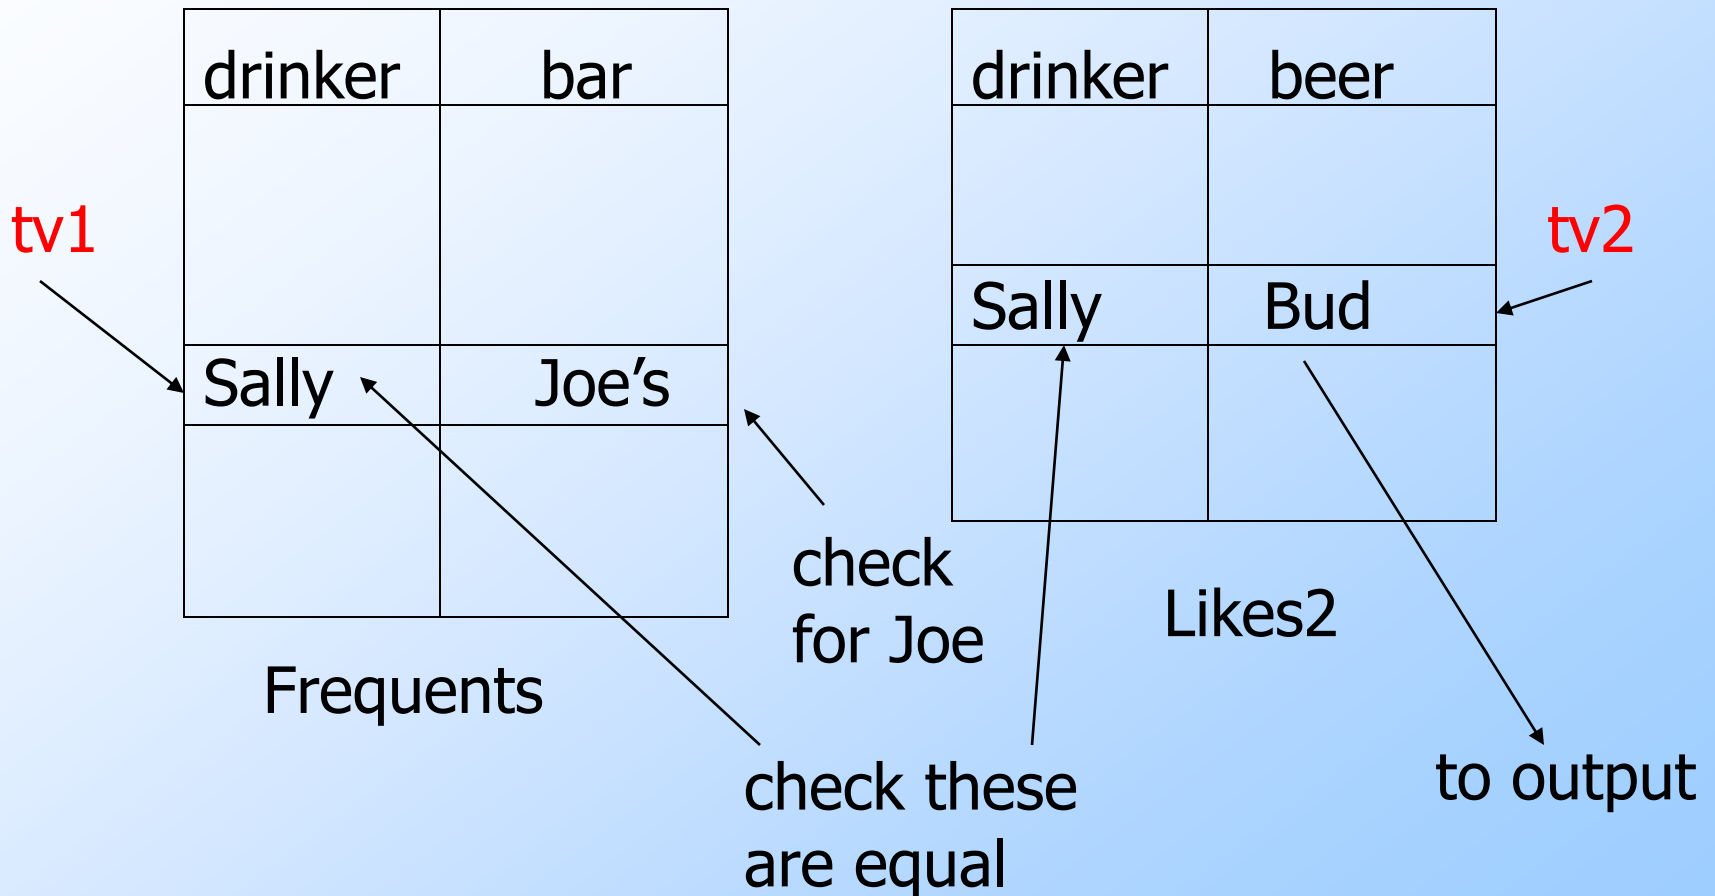
- Almost the same as for single-relation queries:
  1. Start with the **product** of all the relations in the **FROM** clause.
  2. Apply the **selection** condition from the **WHERE** clause.
  3. **Project** onto the list of attributes and expressions in the **SELECT** clause.

# Operational Semantics

- Imagine **one tuple-variable for each relation** in the FROM clause.
  - These tuple-variables visit each combination of tuples, one from each relation.
- If the tuple-variables are pointing to tuples that satisfy the WHERE clause, send these tuples to the SELECT clause.



# Example



# Explicit Tuple-Variables

- Sometimes, a query needs to use two copies of the same relation.
- Distinguish copies by following the relation name by the name of a **tuple-variable**, in the **FROM** clause.
- It's always an option to rename relations this way, even when not essential.

# Example: Self-Join

- From **Beers(name, manf)**, find all pairs of beers by the same manufacturer.
  - Do not produce pairs like (Bud, Bud).
  - Produce pairs in alphabetic order, e.g. (Bud, Miller), not (Miller, Bud).

```
SELECT b1.name, b2.name
FROM Beers b1, Beers b2
WHERE b1.manf = b2.manf AND
      b1.name < b2.name;
```

# Subqueries

- A parenthesized SELECT-FROM-WHERE statement (*subquery*) can be used as a value in a number of places, including **FROM** and **WHERE clauses**.
- **Example**: in place of a relation in the FROM clause, we can use a subquery and then query its result.
  - Must use a tuple-variable to name tuples of the result.


# Example: Subquery in FROM

- Find the beers liked by at least one person who frequents Joe's Bar.

```
SELECT beer
```

```
FROM Likes2, (SELECT drinker  
FROM Frequents  
WHERE bar = 'Joes Bar') JD
```

Drinkers who  
frequent Joe's Bar



```
WHERE Likes2.drinker=JD.drinker;
```

# Subqueries That Return One Tuple

- If a subquery is guaranteed to produce one tuple, then the subquery can be used as a value.
  - Usually, the tuple has one component.
  - A run-time error occurs if there is no tuple or more than one tuple.

# Example: Single-Tuple Subquery

- Using `Sells(bar, beer, price)`, find the bars that serve Miller for the same price Joe charges for Bud.
- Two queries would surely work:
  1. Find the price Joe charges for Bud.
  2. Find the bars that serve Miller at that price.

# Query + Subquery Solution

SELECT bar

FROM Sells

WHERE beer = 'Miller' AND

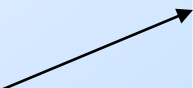
price = (SELECT price

FROM Sells

WHERE bar = 'Joes Bar'

AND beer = 'Bud');

The price at  
which Joe  
sells Bud





# The IN Operator

- `<tuple> IN (<subquery>)` is true if and only if the tuple is a member of the relation produced by the subquery.
  - Opposite: `<tuple> NOT IN (<subquery>)`.
- IN-expressions can appear in WHERE clauses.

## Example: IN

- Using `Beers(name, manf)` and `Likes2(drinker, beer)`, find the name and manufacturer of each beer that Fred likes.

`SELECT *`

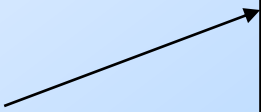
`FROM Beers`

`WHERE name IN (SELECT beer`

`FROM Likes2`

`WHERE drinker = 'Fred');`

The set of  
beers Fred  
likes



# What is the difference?

```
SELECT a  
FROM R, S  
WHERE R.b = S.b;
```

```
SELECT a  
FROM R  
WHERE b IN (SELECT b FROM S);
```

# IN is a Predicate About R's Tuples

```
SELECT a
FROM R
WHERE b IN (SELECT b FROM S);
```

Two 2's

One loop, over  
the tuples of R

| a | b |
|---|---|
| 1 | 2 |
| 3 | 4 |

R

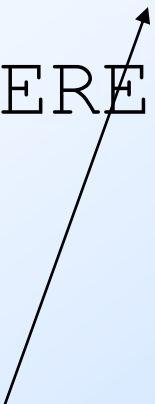
| b | c |
|---|---|
| 2 | 5 |
| 2 | 6 |

S

(1,2) satisfies  
the condition;  
1 is output once.

# This Query Pairs Tuples from R, S

```
SELECT a
FROM R, S
WHERE R.b = S.b;
```



**Double loop**, over  
the tuples of R and S

| a | b |
|---|---|
| 1 | 2 |
| 3 | 4 |

R

| b | c |
|---|---|
| 2 | 5 |
| 2 | 6 |

S

(1,2) with (2,5)  
and (1,2) with  
(2,6) both satisfy  
the condition;  
**1 is output twice.**

# The Exists Operator

- **EXISTS**(<subquery>) is true if and only if the subquery result is not empty.
- **Example**: From **Beers(name, manf)** , find those beers that are the unique beer by their manufacturer.

# Example: EXISTS

```
SELECT name  
FROM Beers b1  
WHERE NOT EXISTS (
```

Notice scope rule: manf refers to closest nested FROM with a relation having that attribute.

Set of beers with the same manf as b1, but not the same beer

```
SELECT *  
FROM Beers  
WHERE manf = b1.manf AND  
      name <> b1.name);
```

Notice the SQL "not equals" operator

# The Operator ANY

- $x = \text{ANY}(\langle \text{subquery} \rangle)$  is a boolean condition that is true iff  $x$  equals at least one tuple in the subquery result.
  - " $=$ " could be any comparison operator.
- Example:  $x \geq \text{ANY}(\langle \text{subquery} \rangle)$  means  $x$  is not the uniquely smallest tuple produced by the subquery.
  - Note tuples must have one component only.



# The Operator ALL

- $x <> \text{ALL}(<\text{subquery}>)$  is true iff for every tuple  $t$  in the relation,  $x$  is not equal to  $t$ .
  - That is,  $x$  is not in the subquery result.
- $<>$  can be any comparison operator.
- **Example:**  $x \geq \text{ALL}(<\text{subquery}>)$  means there is no tuple larger than  $x$  in the subquery result.

# Example: ALL

- From **Sells(bar, beer, price)**, find the beer(s) sold for the highest price.

SELECT beer

FROM Sells

WHERE price >= ALL(  
SELECT price  
FROM Sells);

price from the outer  
Sells must not be  
less than any price.

# Union, Intersection, and Difference

□ Union, intersection, and difference of relations are expressed by the following forms, each involving subqueries:

□ (<subquery>) UNION (<subquery>)

□ (<subquery>) INTERSECT (<subquery>)

□ (<subquery>) MINUS (<subquery>)

/\* Some DBMSs use EXCEPT instead of MINUS \*/

# Example: Intersection

- Using `Likes2(drinker, beer)`, `Sells(bar, beer, price)`, and `Frequents(drinker, bar)`, find the drinkers and beers such that:
  1. The drinker likes the beer, and
  2. The drinker frequents at least one bar that sells the beer.

Notice trick:  
subquery is  
really a stored  
table.

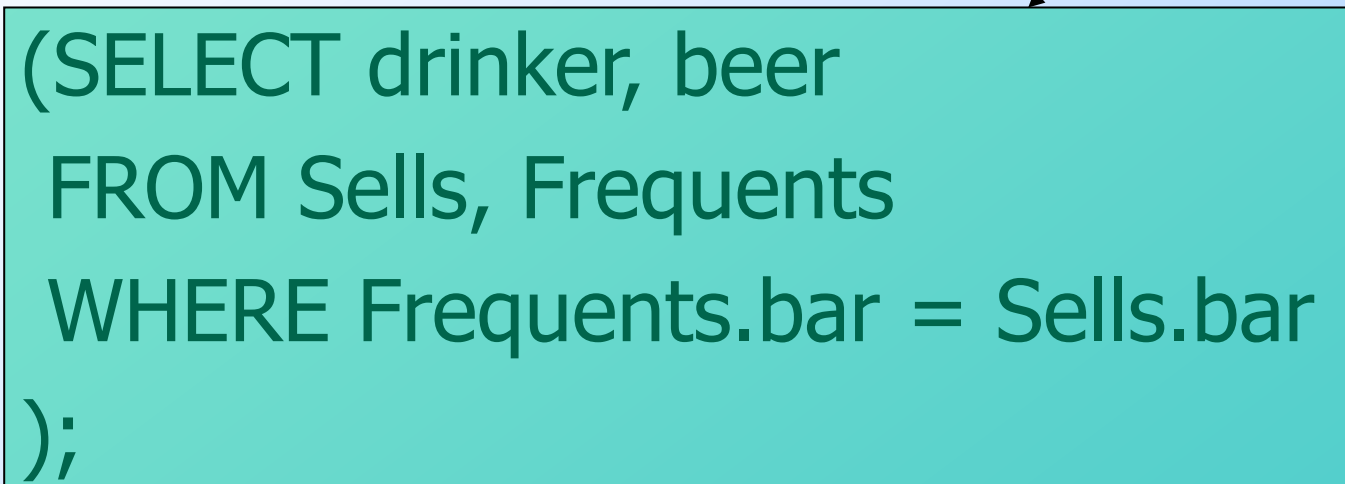
# Solution



```
(SELECT * FROM Likes2)
```

INTERSECT

The drinker frequents  
a bar that sells the  
beer.



```
(SELECT drinker, beer  
FROM Sells, Frequents  
WHERE Frequents.bar = Sells.bar  
);
```

# Bag Semantics

- Although the SELECT-FROM-WHERE statement uses bag semantics, the default for union, intersection, and difference is set semantics.
- That is, duplicates are eliminated as the operation is applied.

# Motivation: Efficiency

- When doing projection, it is easier to avoid eliminating duplicates.
  - Just work tuple-at-a-time.
- For **intersection** or **difference**, it is **most efficient to sort** the relations **first**.
  - At that point you may as well eliminate the duplicates anyway.

# Controlling Duplicate Elimination

- Force the result to be a set by  
SELECT **DISTINCT** . . .
- Force the result to be a bag (i.e., don't eliminate duplicates) by ALL, as in  
. . . **UNION ALL** . . .



# Example: DISTINCT

- From `Sells(bar, beer, price)`, find all the different prices charged for beers:

```
SELECT DISTINCT price  
FROM Sells;
```

- Notice that without `DISTINCT`, each price would be listed as many times as there were bar/beer pairs at that price.

# Example: ALL

- Using relations **Frequents(drinker, bar)** and **Likes2(drinker, beer)**:

```
(SELECT drinker FROM Frequents)
  MINUS ALL /* Oracle doesn't support */
(SELECT drinker FROM Likes2);
```

- Lists drinkers who frequent more bars than they like beers, and does so as many times as the difference of those counts.

# Join Expressions

- SQL provides several versions of (bag) joins.
- These expressions can be stand-alone queries or used in place of relations in a FROM clause.

# Products and Natural Joins

- Natural join:

R **NATURAL JOIN** S;

- Product:

R **CROSS JOIN** S;

- **Example:**

Likes2 **NATURAL JOIN** Sells;

- Relations can be **parenthesized subqueries**, as well.

# Theta Join

- $R \text{ JOIN } S \text{ ON } \langle \text{condition} \rangle$
- **Example:** using  $\text{Drinkers}(\text{name}, \text{addr})$  and  $\text{Frequents}(\text{drinker}, \text{bar})$ :

```
Drinkers JOIN Frequents ON  
    name = drinker;
```

gives us all  $(d, a, d, b)$  quadruples such that drinker  $d$  lives at address  $a$  and frequents bar  $b$ .